

Less State, More Logs!

Scaling Ethereum With A Log-Based Storage Architecture

Zsolt Felföldi

Ethereum Foundation (GETH)

zsselfoldi@ethereum.org

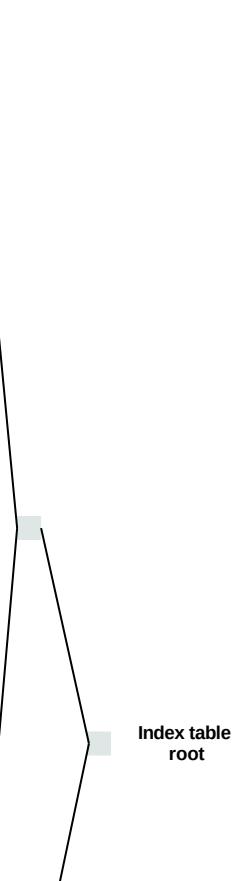
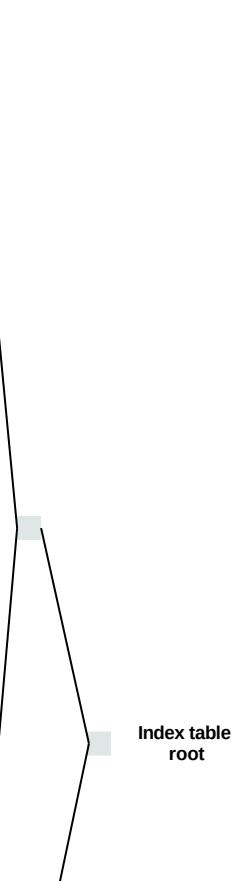
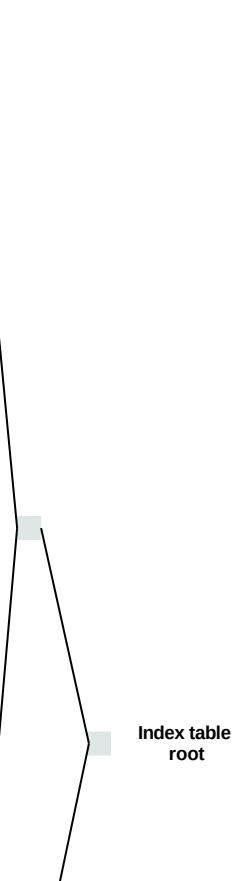
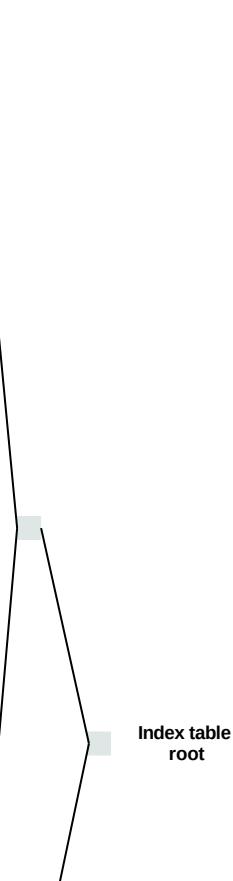
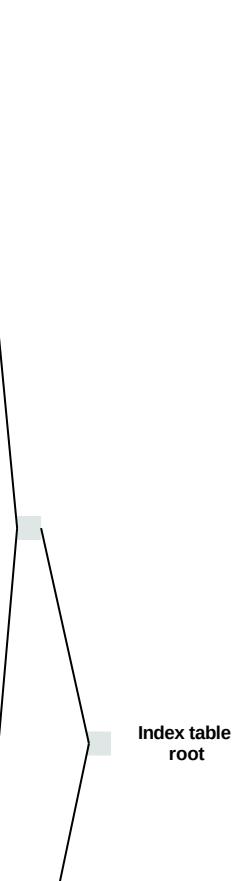
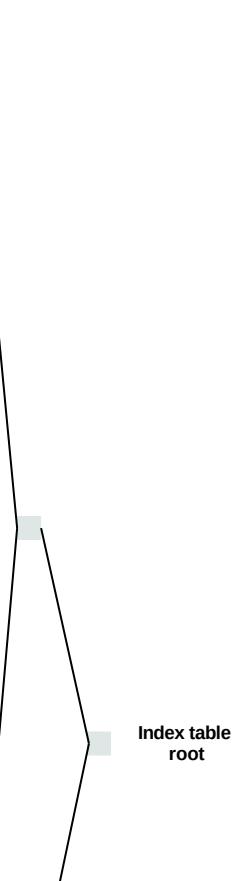
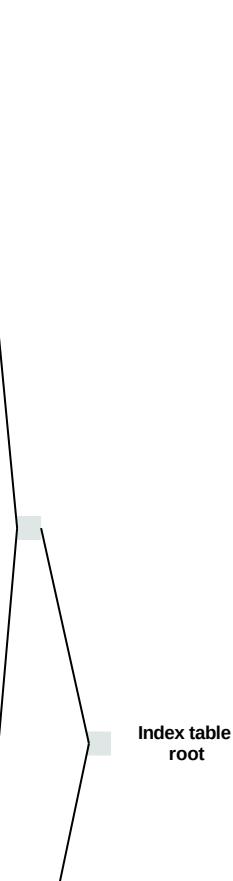
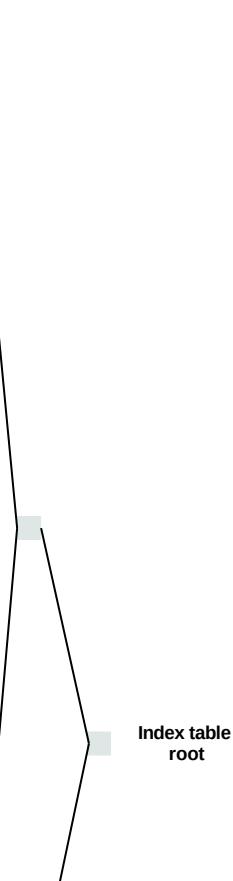
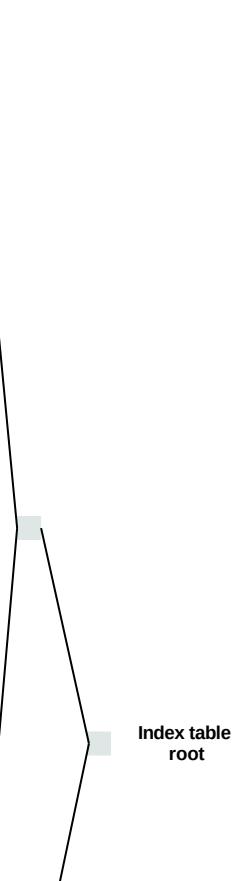
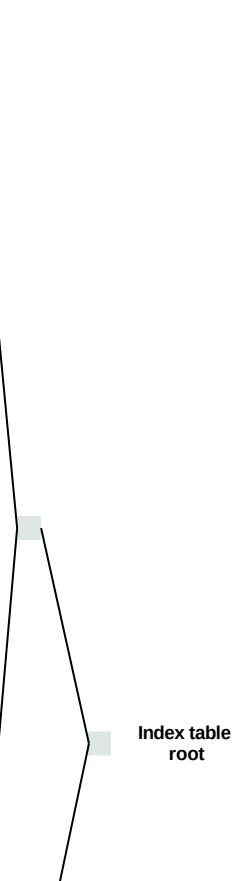
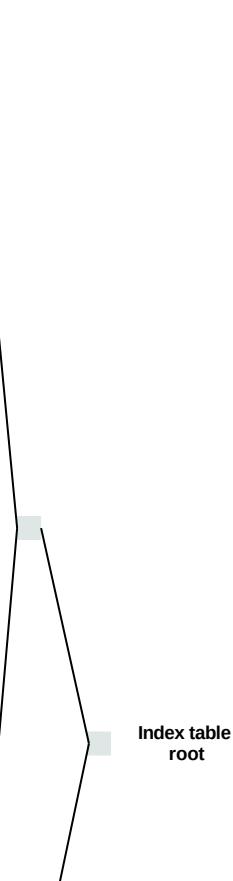
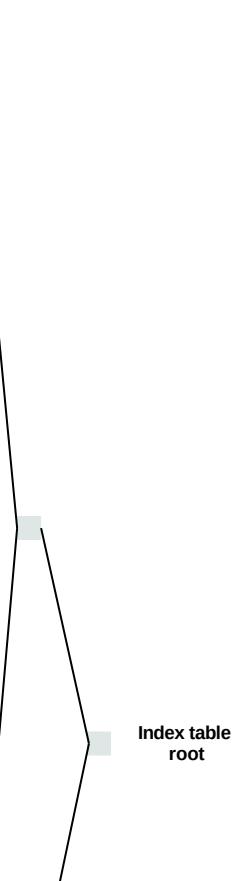
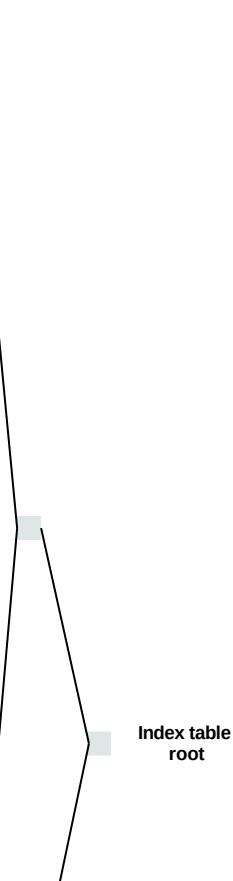
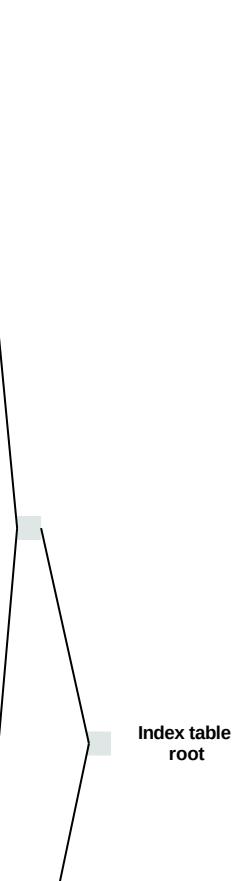
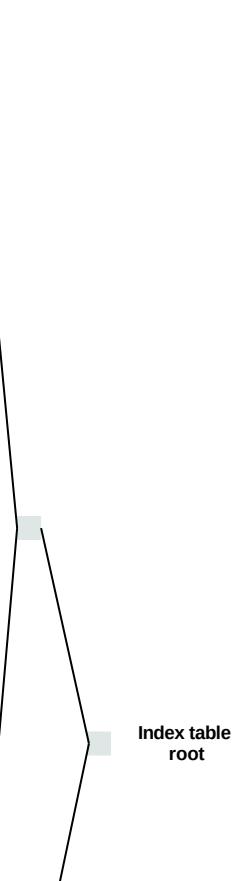
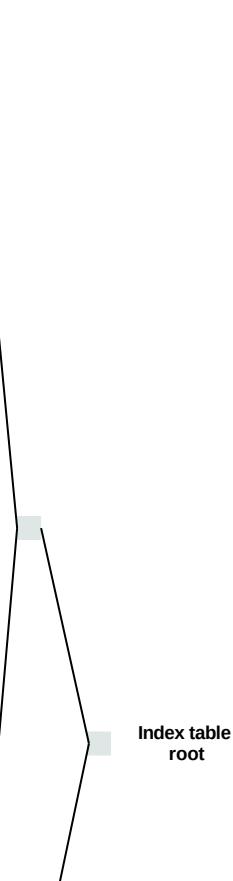
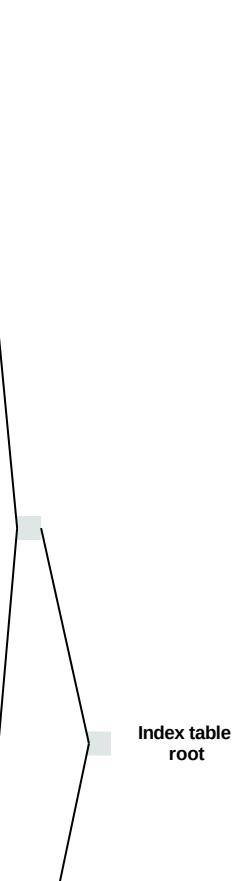
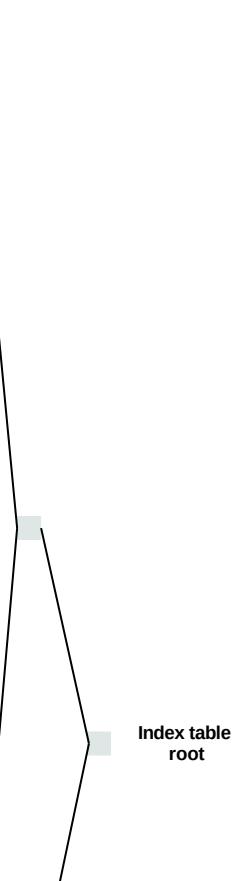
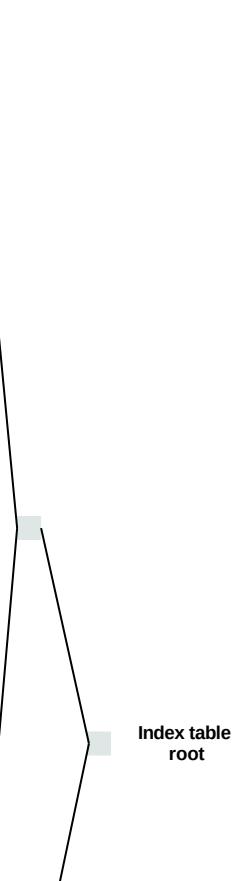
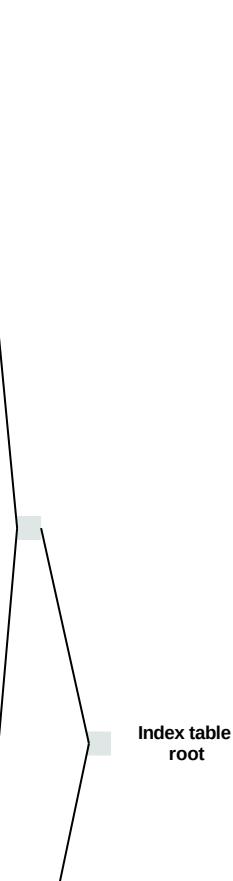
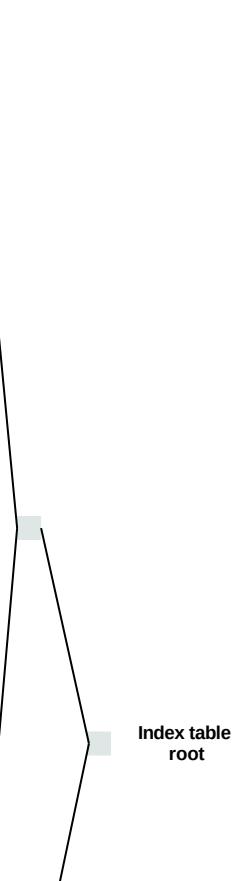
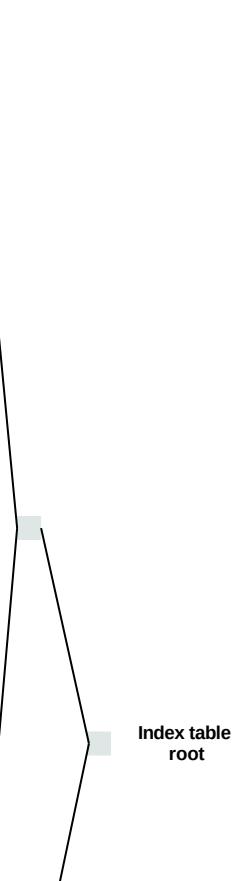
Trustless Execution Layer API

JSON RPC	Trustless execution layer REST API
<code>eth_blockNumber</code>	none (assumed to be known by the client)
<code>eth_getBalance</code>	<code>/eth/v1/exec/state</code>
<code>eth_getBlockByHash</code>	<code>/eth/v1/exec/blocks</code> if canonicalness needs to be proven, <code>/eth/v1/exec/history</code> * *alternative: <code>/eth/v1/exec/headers</code> from queried block to head
<code>eth_getBlockByNumber</code>	<code>/eth/v1/exec/blocks</code> , <code>/eth/v1/exec/history</code> * *alternative: <code>/eth/v1/exec/headers</code> from queried block to head
<code>eth_getBlockReceipts</code>	<code>/eth/v1/exec/block_receipts</code>
<code>eth_getBlockTransactionCountByHash</code>	<code>/eth/v1/exec/transaction_by_index</code>
<code>eth_getBlockTransactionCountByNumber</code>	<code>/eth/v1/exec/history</code> * or <code>/eth/v1/exec/headers</code> from queried block to head <code>/eth/v1/exec/transaction_by_index</code>
<code>eth_getCode</code>	<code>/eth/v1/exec/state</code>
<code>eth_getLogs</code>	<code>/eth/v1/exec/logs</code> * *alternative: <code>/eth/v1/exec/headers</code> from range start to head, <code>/eth/v1/exec/block_receipts</code> for all bloom filter matches
<code>eth_getProof</code>	<code>/eth/v1/exec/state</code>
<code>eth_getStorageAt</code>	<code>/eth/v1/exec/state</code>
<code>eth_getTransactionByBlockHashAndIndex</code>	<code>/eth/v1/exec/headers</code> , <code>/eth/v1/exec/transaction_by_index</code>
<code>eth_getTransactionByBlockNumberAndIndex</code>	<code>/eth/v1/exec/history</code> *, <code>/eth/v1/exec/headers</code> for queried block, <code>/eth/v1/exec/transaction_by_index</code> *alternative: <code>/eth/v1/exec/headers</code> from queried block to head
<code>eth_getTransactionByHash</code>	<code>/eth/v1/exec/transaction</code> <code>/eth/v1/exec/transaction_position</code> * if proof of canonical inclusion position is needed *alternative for canonical transactions: <code>/eth/v1/exec/headers</code> from inclusion block to head <code>/eth/v1/exec/transaction_by_index</code>

- Endpoint format similar to the beacon chain REST API
- Works together with a beacon light client; answers all queries with Merkle proofs based on latest execution block hash
- New functionality based on the Trustless Log Index (EIP-7745B)

See full draft proposal at:

<https://github.com/zsfelfoldi/tli>

Type	Value	Block #	Tx #	Log #	
0 (Block)	0x42f66a2e9f9c68... (block hash)	41			
0 (Block)	0x66f42ef12b140e... (block hash)	42			
0 (Block)	0x978ce0036b6d1c... (block hash)	43			
0 (Block)	0xbff98e6cb26f6ff... (block hash)	40			
1 (Tx)	0x7046035b326ab2... (tx hash)	43	0		
1 (Tx)	0xa75590c9ced728... (tx hash)	42	1		
1 (Tx)	0xca2d12d1b8132d... (tx hash)	42	0		
2 (Address)	0x00..00c02aaa39... (WETH)	42	0	0	
2 (Address)	0x00..00c02aaa39... (WETH)	42	1	0	
2 (Address)	0x00..00dac17f95... (USDT)	42	0	1	
2 (Address)	0x00..00dac17f95... (USDT)	43	0	0	
3 (Topic0)	0x7fcf532c15f0a6... (Withdrawal)	42	1	0	
3 (Topic0)	0xddf252ad1be2c8... (Transfer)	42	0	0	
3 (Topic0)	0xddf252ad1be2c8... (Transfer)	42	0	1	
3 (Topic0)	0xddf252ad1be2c8... (Transfer)	43	0	0	
4 (Topic1)	0x00..004a5e9a3b... (Alice)	42	0	0	
4 (Topic1)	0x00..004a5e9a3b... (Alice)	43	0	0	
4 (Topic1)	0x00..00f2a5fd4d... (Bob)	42	0	1	
4 (Topic1)	0x00..00f2a5fd4d... (Bob)	42	1	0	
5 (Topic2)	0x00..004a5e9a3b... (Alice)	42	0	1	
5 (Topic2)	0x00..00ac8cbe20... (Carol)	43	0	0	
5 (Topic2)	0x00..00f2a5fd4d... (Bob)	42	0	0	

 Tree node

 Index table root

 SSZ List count (22)

The Trustless Log Index

(introducing index tables)

- A lexicographically ordered list of entries in a given block range
 - Log addresses and topics
 - Block hashes, transaction hashes
- Stored with position information
- Hashed into a Binary Merkle Tree
- Inclusion and exclusion proofs for specific entries or search patterns

Type	Value	Block #	Tx #	Log #	
0 (Block)	0x42f66a2e9f9c68... (block hash)	41			
0 (Block)	0x66f42ef12b140e... (block hash)	42			
0 (Block)	0x978ce0036b6d1c... (block hash)	43			
0 (Block)	0xbff98e6cb26f6ff... (block hash)	40			
1 (Tx)	0x7046035b326ab2... (tx hash)	43	0		
1 (Tx)	0xa75590c9ced728... (tx hash)	42	1		
1 (Tx)	0xca2d12d1b8132d... (tx hash)	42	0		
2 (Address)	0x00..00c02aaa39... (WETH)	42	0	0	
2 (Address)	0x00..00c02aaa39... (WETH)	42	1	0	
2 (Address)	0x00..00dac17f95... (USDT)	42	0	1	
2 (Address)	0x00..00dac17f95... (USDT)	43	0	0	
3 (Topic0)	0x7fcf532c15f0a6... (Withdrawal)	42	1	0	
3 (Topic0)	0xddf252ad1be2c8... (Transfer)	42	0	0	
3 (Topic0)	0xddf252ad1be2c8... (Transfer)	42	0	1	
3 (Topic0)	0xddf252ad1be2c8... (Transfer)	43	0	0	
4 (Topic1)	0x00..004a5e9a3b... (Alice)	42	0	0	
4 (Topic1)	0x00..004a5e9a3b... (Alice)	43	0	0	
4 (Topic1)	0x00..00f2a5fd4d... (Bob)	42	0	1	
4 (Topic1)	0x00..00f2a5fd4d... (Bob)	42	1	0	
5 (Topic2)	0x00..004a5e9a3b... (Alice)	42	0	1	
5 (Topic2)	0x00..00ac8cbe20... (Carol)	43	0	0	
5 (Topic2)	0x00..00f2a5fd4d... (Bob)	42	0	0	

Tree node (not relevant for the proof)

Proof node (included in the proof)

Hash recalculated by verifier

SSZ List count (22)

Proof example: exclusion proof

Block range: #40..#43

- event type: 1 (transaction hash)
- index value: OXdeadbeefcafe...

Results: not found

- Prover looks up the searched entry in the table using binary search
- Searched entry does not exist, finds two adjacent entries
- Creates a multi-proof proving those two entries
- Proof cost per index table: two entries (cca 40-45 bytes per entry with compact encoding) plus $\text{ceil}(\log_2(\text{count}))$ 32 byte proof nodes

Type	Value	Block #	Tx #	Log #
0 (Block)	0x42f66a2e9f9c68... (block hash)	41		
0 (Block)	0x66f42ef12b140e... (block hash)	42		
0 (Block)	0x978ce0036b6d1c... (block hash)	43		
0 (Block)	0xbff98e6cb26f6ff... (block hash)	40		
1 (Tx)	0x7046035b326ab2... (tx hash)	43	0	
1 (Tx)	0xa75590c9ced728... (tx hash)	42	1	
1 (Tx)	0xca2d12d1b8132d... (tx hash)	42	0	
2 (Address)	0x00..00c02aaa39... (WETH)	42	0	0
2 (Address)	0x00..00c02aaa39... (WETH)	42	1	0
2 (Address)	0x00..00dac17f95... (USDT)	42	0	1
2 (Address)	0x00..00dac17f95... (USDT)	43	0	0
3 (Topic0)	0x7fcf532c15f0a6... (Withdrawal)	42	1	0
3 (Topic0)	0xddf252ad1be2c8... (Transfer)	42	0	0
3 (Topic0)	0xddf252ad1be2c8... (Transfer)	42	0	1
3 (Topic0)	0xddf252ad1be2c8... (Transfer)	43	0	0
4 (Topic1)	0x00..004a5e9a3b... (Alice)	42	0	0
4 (Topic1)	0x00..004a5e9a3b... (Alice)	43	0	0
4 (Topic1)	0x00..00f2a5fd4d... (Bob)	42	0	1
4 (Topic1)	0x00..00f2a5fd4d... (Bob)	42	1	0
5 (Topic2)	0x00..004a5e9a3b... (Alice)	42	0	1
5 (Topic2)	0x00..00ac8cbe20... (Carol)	43	0	0
5 (Topic2)	0x00..00f2a5fd4d... (Bob)	42	0	0

Tree node (not relevant for the proof)

Proof node (included in the proof)

Hash recalculated by verifier

SSZ List count (22)

Proof example: log pattern search

Block range: #40..#43

Address: 0x00..00dac17f95... (USDT)

Topic0: 0xddf252ad1be2c8... (Transfer)

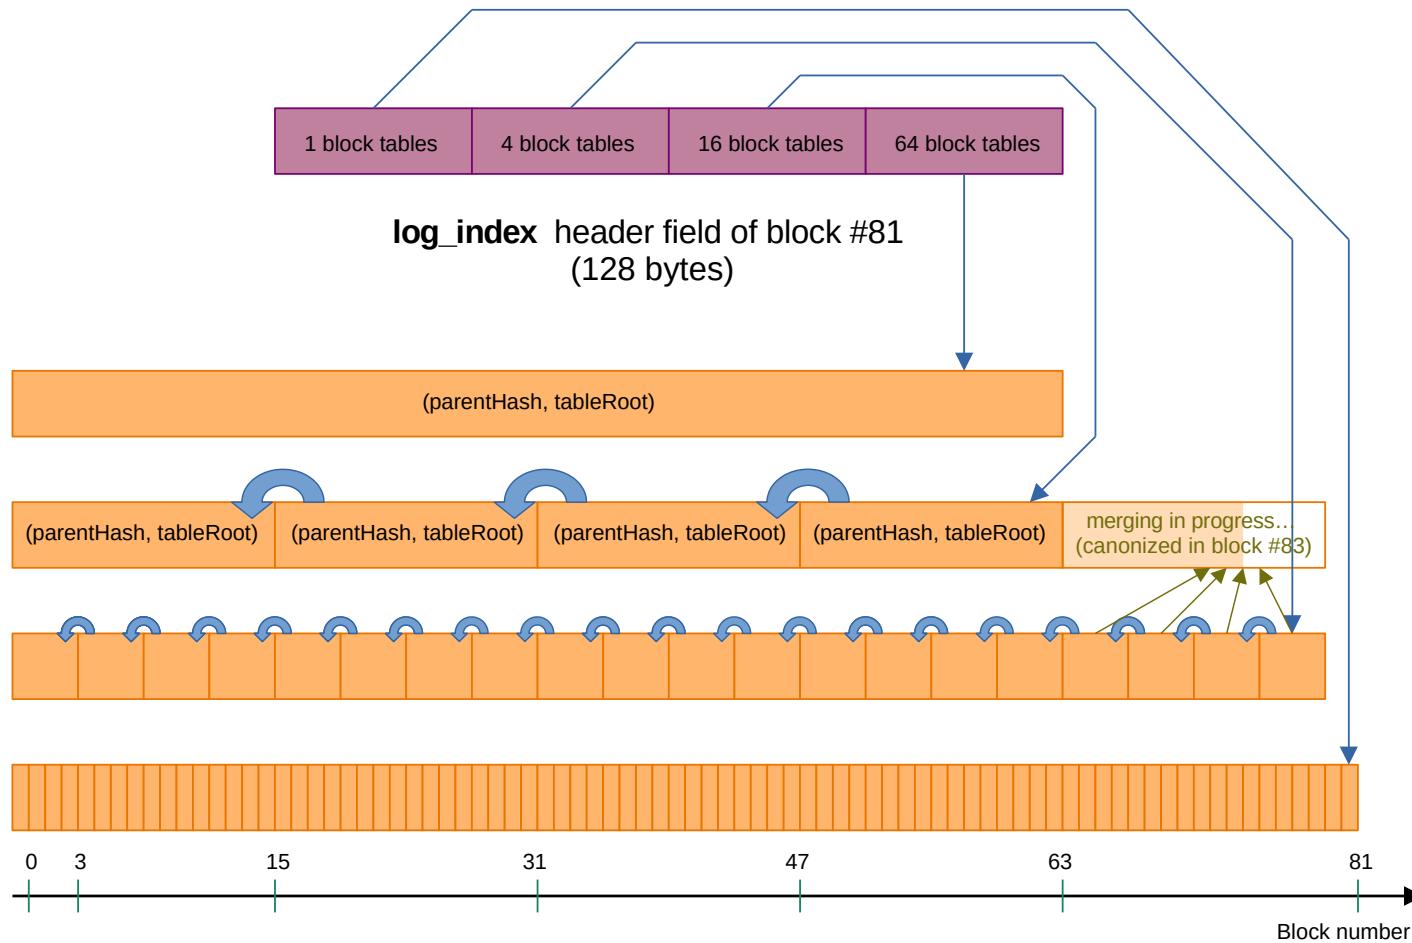
Topic1: 0x00..00f2a5fd4d... (Bob)

Topic2: anything

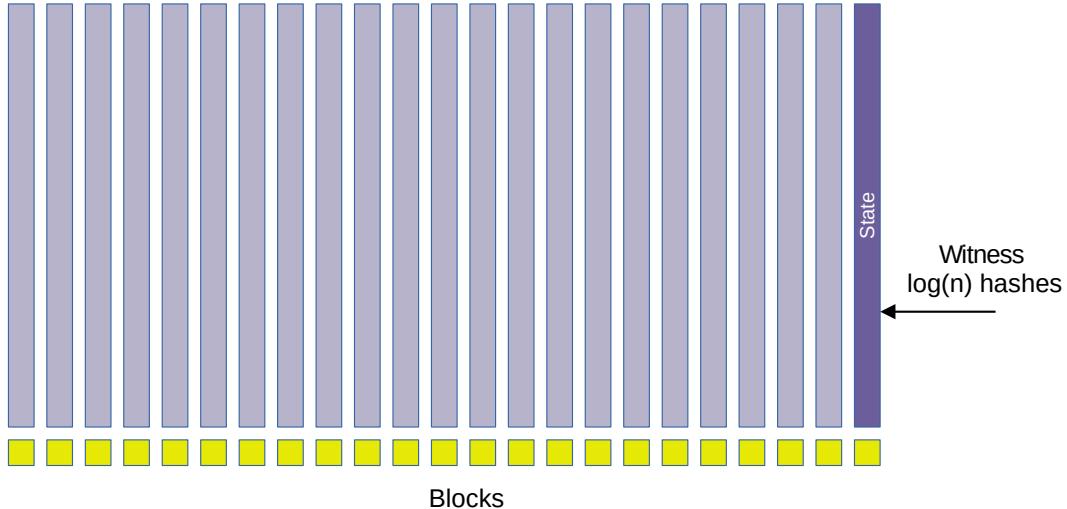
Results: [{Block #42 Tx #0 Log #1}]

- All occurrences of **Topic1 = Bob** are proven:
 - {Block #42 Tx #0 Log #1}
 - {Block #42 Tx #1 Log #0}
- **Topic0 = Withdrawal** in {Block #42 Tx #1 Log #0} (*no match*)
- **Block #42** root hash is proven
 - Receipt proves that {Block #42 Tx #0 Log #1} is a *match*

Covering the chain history with table chains



Logs vs State

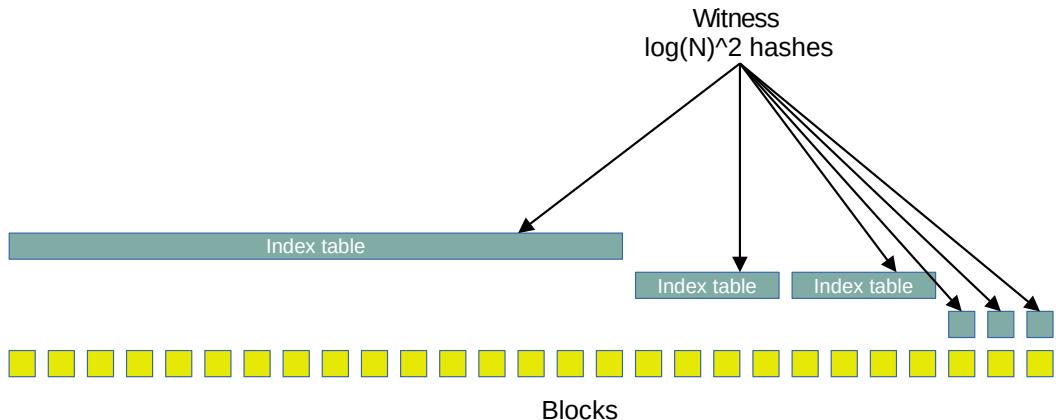


Pro state tree:

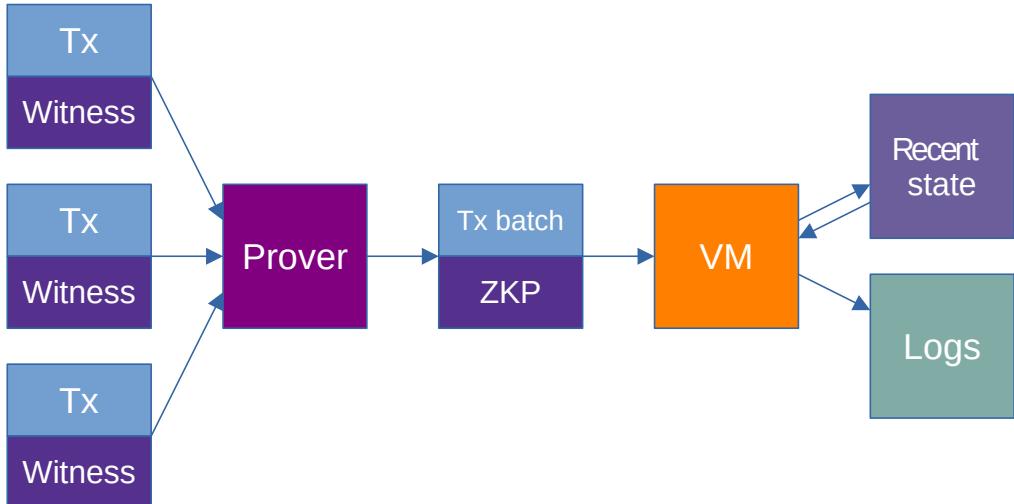
- Smaller witness size

Pro index tables:

- Very cheap to add new entries
- Most of the processing is async/off-chain
- Tables are not changed, only merged (big ones are unchanged for a long time)
- Easy to distribute, no witness provider centralization issue
- Time information is available



Off-chain authorization, pre-checks and witness processing

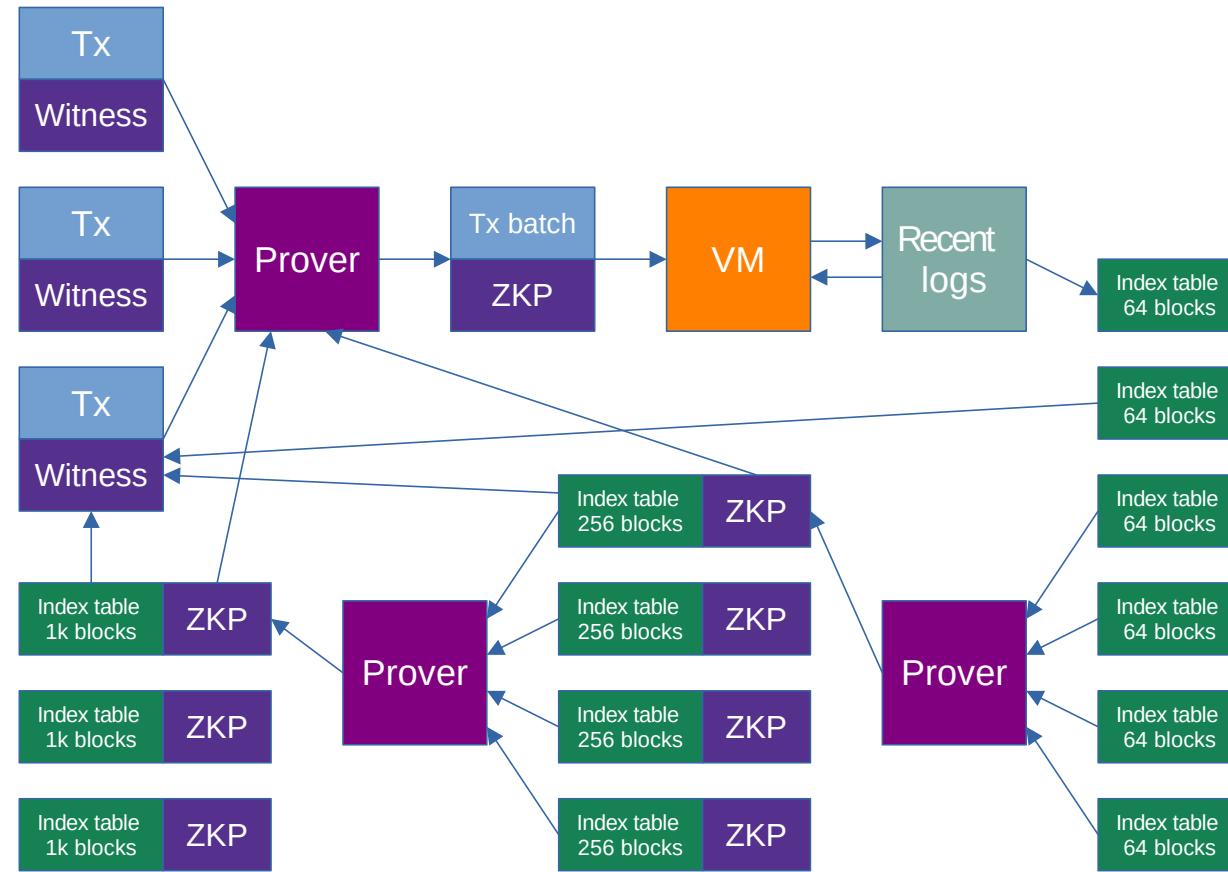


- Witnesses, signatures and other pre-check data can be processed in a ZKP
- Proving cost settled with provers on private channels, on-chain proof verification cost only
- Cheap pre-checks good for
 - Full or partial statelessness
 - Account abstraction, big PQ signatures

Logs as contract storage

- Implement the Trustless Log Index (EIP-7745B, much simpler than original EIP-7745)
 - Very cheap to append, easy to expire, easy to distribute
 - Big index tables merged with ZK proofs (immutable once max table size reached)
 - Add erasure coding, ensure data availability until tables expired
 - Keeping relevant witnesses from expired tables is the user's responsibility
- Make logs accessible from the EVM
 - Recent ones without witness, older ones with witness
- Can simulate state but more flexible
- Append-only, content/time addressed storage
 - Expiry can be controlled by contract (only search non-expired index range)
- Parallel execution friendly (sets, lists, tables without write collisions)
- Multi-chain friendly (easy to list relevant events since last update)

Prover infrastructure of a log-based storage architecture



Thank you for your attention!

Trustless Log Index repository:

<https://github.com/zsfelfoldi/tli>

Contact:

[zsfelfoldi @ discord](#)

zsfelfoldi@ethereum.org